

# Communication in Distributed Systems – Fundamental Concepts

---

.....

# Lecture Overview

---

- Understanding basic terminologies in communication in distributed systems
- Understanding key concepts in communication in distributed systems

# Outline

---

- Communication entities, paradigm, roles/responsibilities
- Key issues in communication in distributed systems
- Protocols
- Processing requests
- Summary

# Learning Material

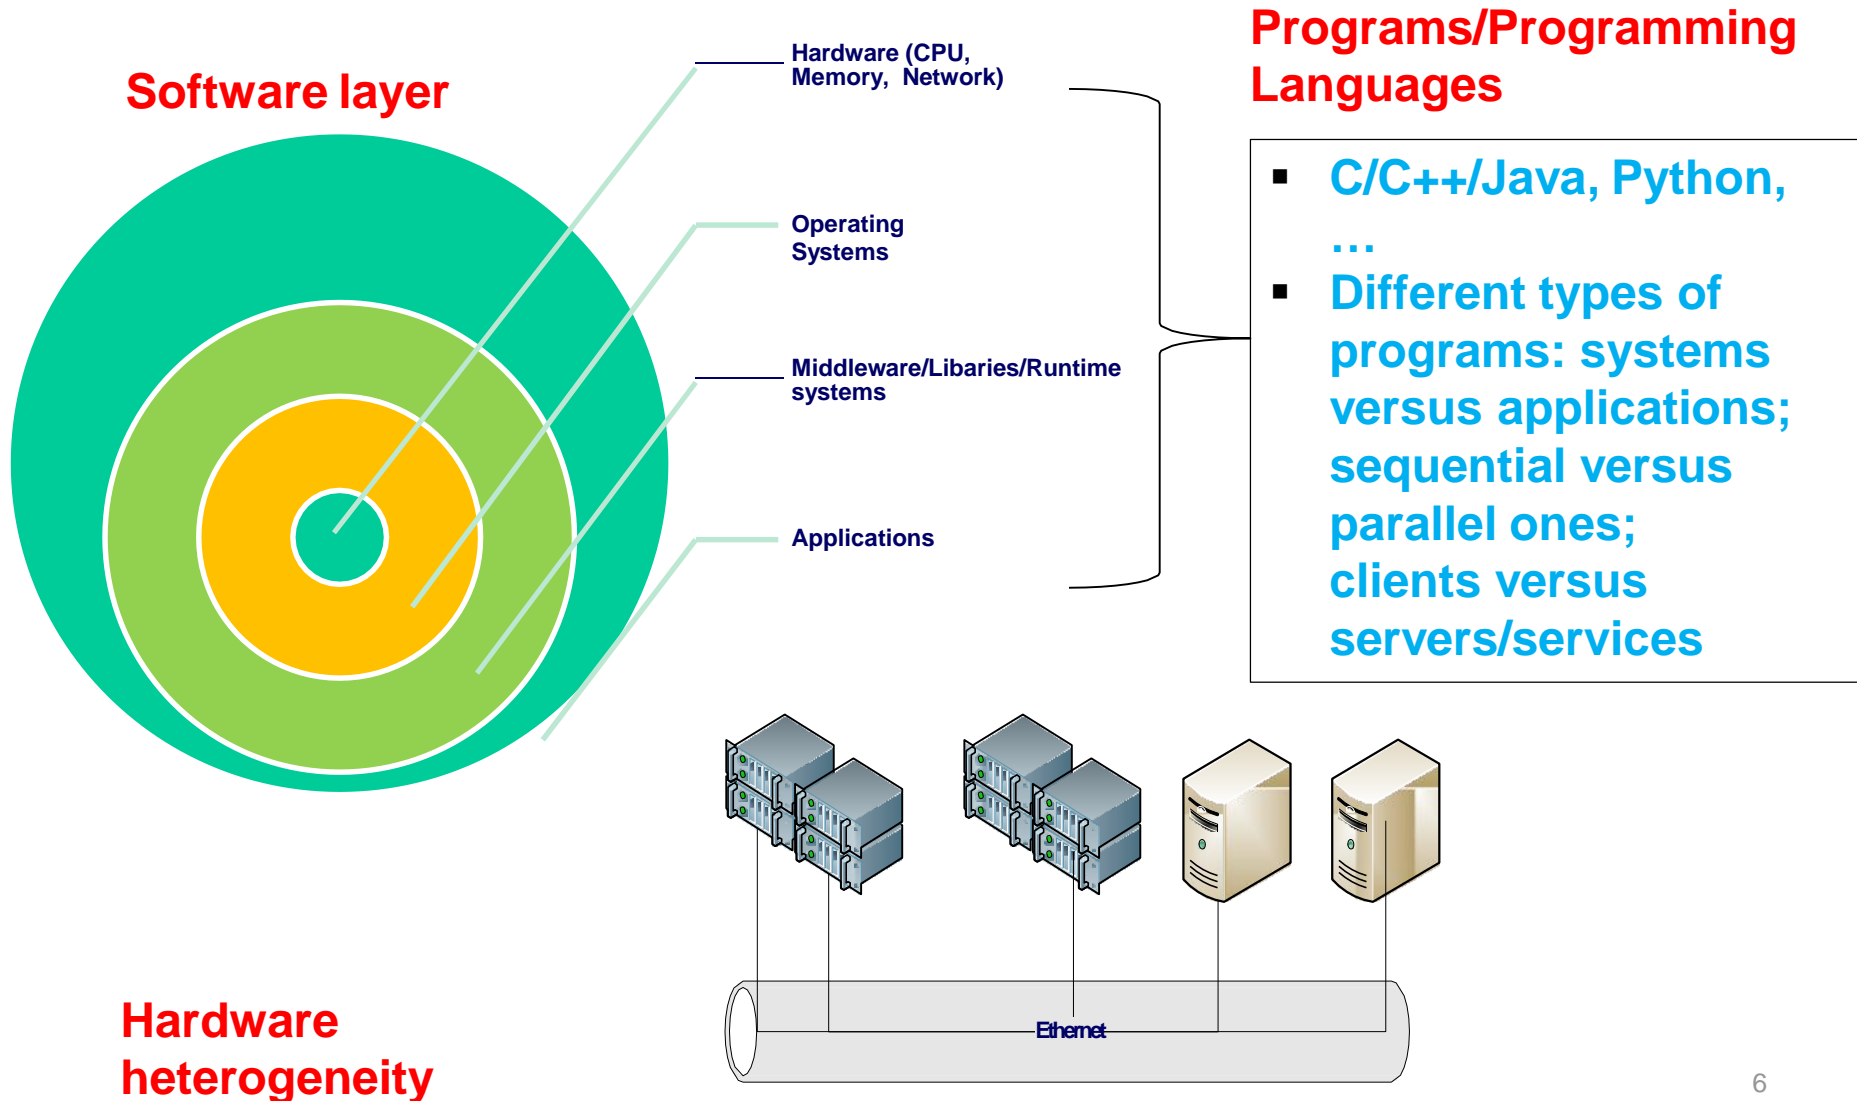
---

- Main reading:
  - George Coulouris, Jean Dollimore, Tim Kindberg, Gordon Blair, „Distributed Systems – Concepts and Design“, 5nd Edition
    - Chapters 2,3, 7.
  - Craig Hunt, TCP/IP Network Administration, 3edition, 2002, O'Reilly.

---


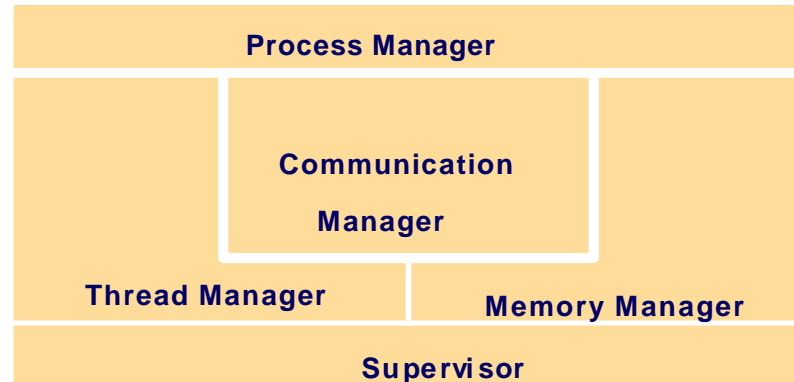
# **COMMUNICATION ENTITIES, PARADIGM, AND ROLES**

# Hardware, software layer, programs



# Systems layers and Core OS functionality

## Core OS functionality



OS: Kernel,  
Libraries & Servers




Platform

Node 1

Node 2


Source: Coulouris, Dollimore,  
Kindberg

# Process vs Thread



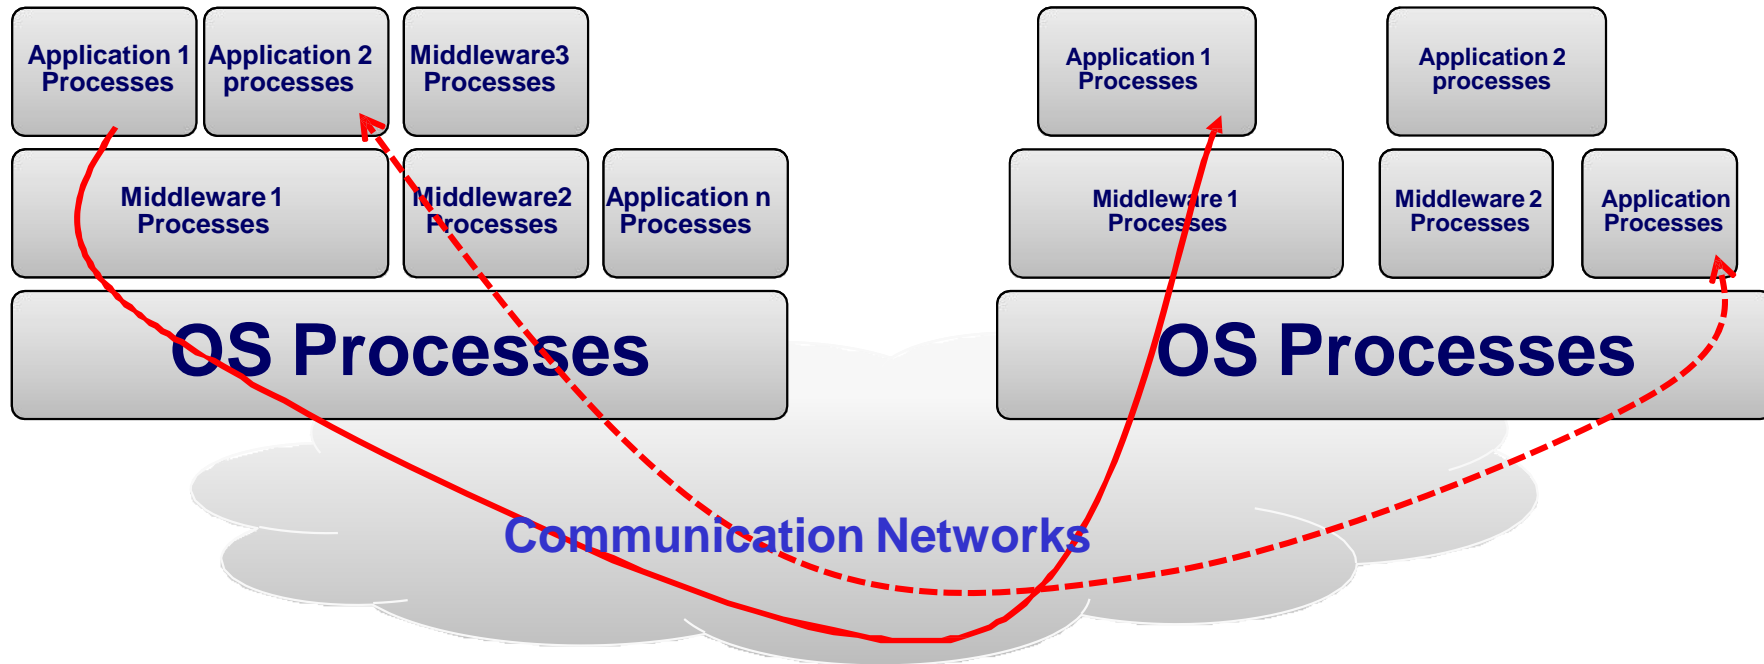
## Within a non distributed OS

- Process – the program being executed by the OS
- Threads within a process
- Switching thread context is much cheaper than that for the process context
- Blocking calls in a thread do not block the whole process






# Communication entities



## Communication in distributed systems

- **between processes within a single** application/middleware/service
- **among processes belonging to different** applications/middleware/services
- **Among computing nodes** which have no concept of processes (e.g. sensors)

# Space and Time in Communication



**Q: why is understanding time and space uncoupling important for implementing communication in distributed systems?**

# Communication networks in Distributed Systems

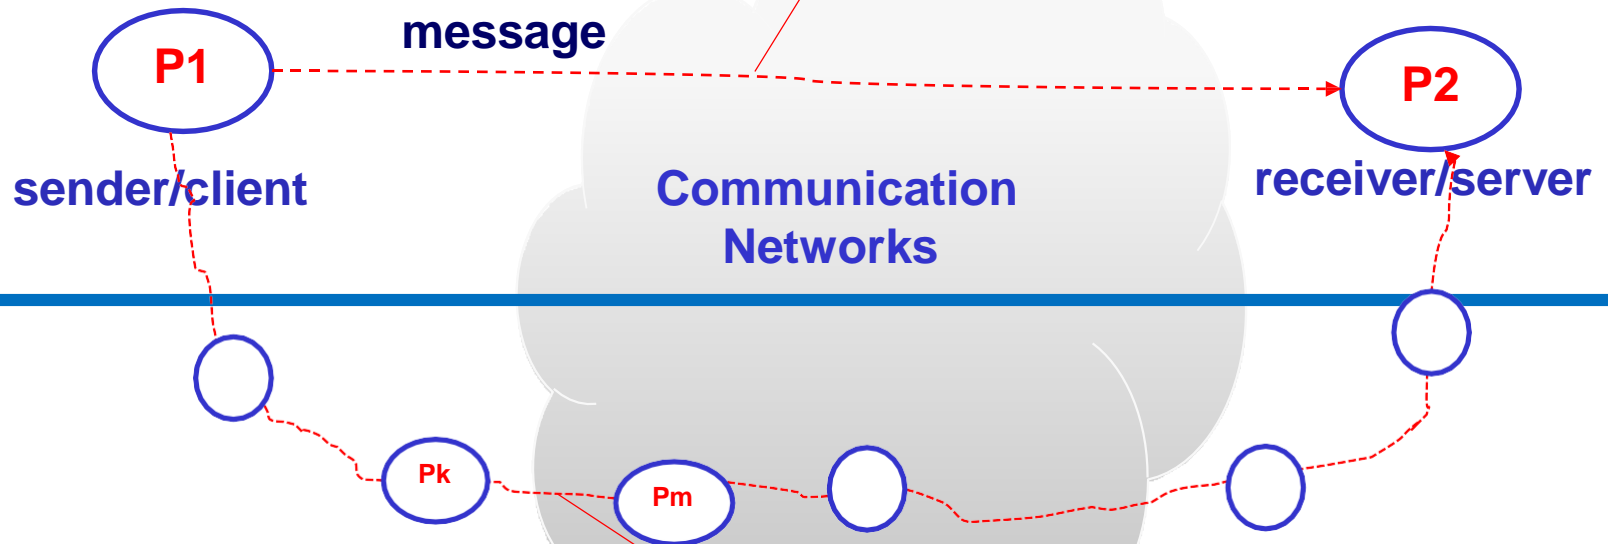
---

- **Maybe designed for specific *types of environments***
  - High performance computing, M2M (Machine-to-Machine), building/home/city management, etc.
  - Events, voices, documents, image data, etc.
- **Distributed, different *network spans***
  - Personal area networks (PANs), local area networks (LANs), campus area networks (CANs), metropolitan area networks (MANs), and wide area networks (WANs)
  - Communication entities are placed in different locations
- **Different *layered networks* for distributed systems**
  - Physical versus overlay network topologies (virtual network topologies atop physical networks)

# Layered communication

In the view of P1 and P2

End-to-end process-to-process communication  
e.g., email [abc@cse.scu.edu](mailto:abc@cse.scu.edu)  
to [ab@gmail.com](mailto:ab@gmail.com)




Holistic system view



End-to-end process-to-process communication

# Communication Patterns

**One-to-one/client-server**



**Group communication**



# Identifiers of entities participating in communication

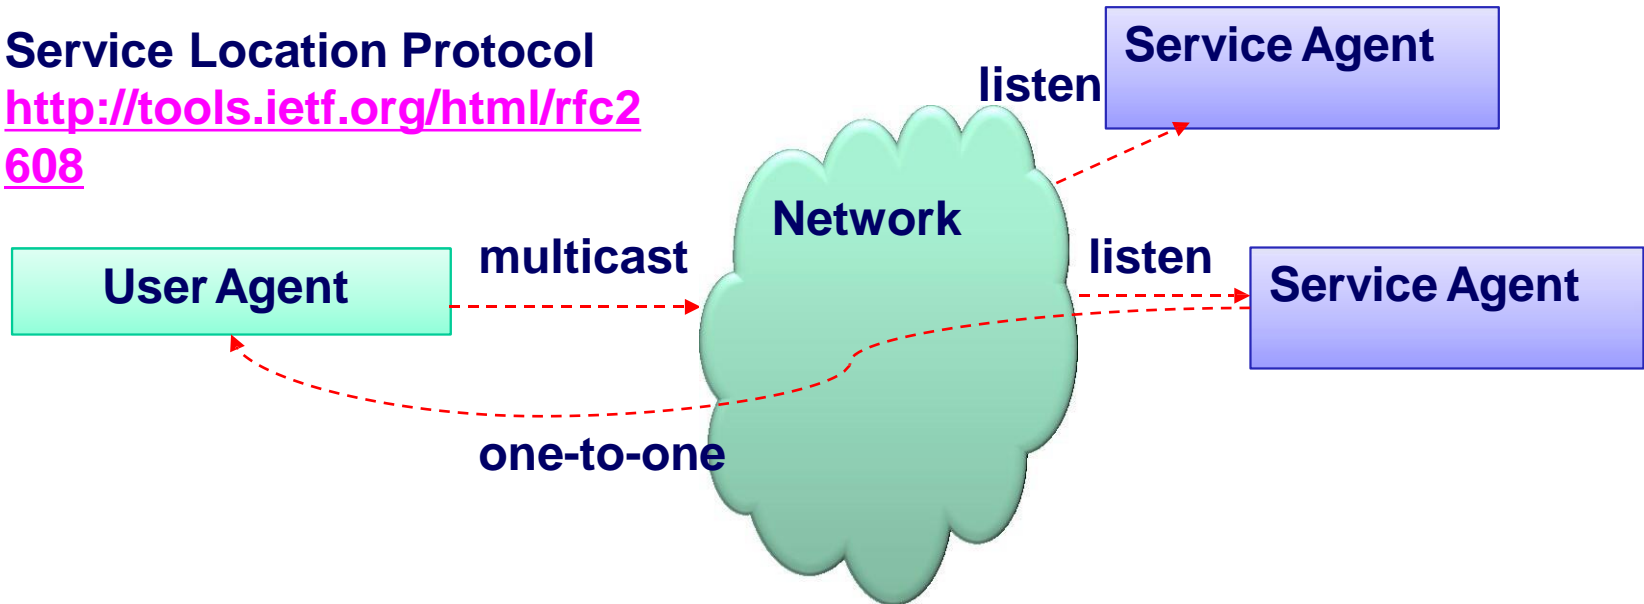
---

- Communication cannot be done without knowing
- identifiers (names) of participating entities
  - Local versus global identifier
  - Individual versus group identifier
- Multiple layers/entities -> different forms of identifiers
  - Process ID in an OS
  - Machine ID: name/IP address
  - Access point: (machine ID, port number)
  - A unique communication ID in a communication network
  - Emails for humans
  - Group ID

# Examples of communication patterns (I)


## Service Location Protocol

<http://tools.ietf.org/html/rfc2608>



- A User Agent wants to find a Service Agent
- Different roles and different communication patterns
- Get <http://jslp.sourceforge.net/> and play samples to see how it works

# Examples of communication patterns (2)



## ■ MPI (Message Passing Interface)

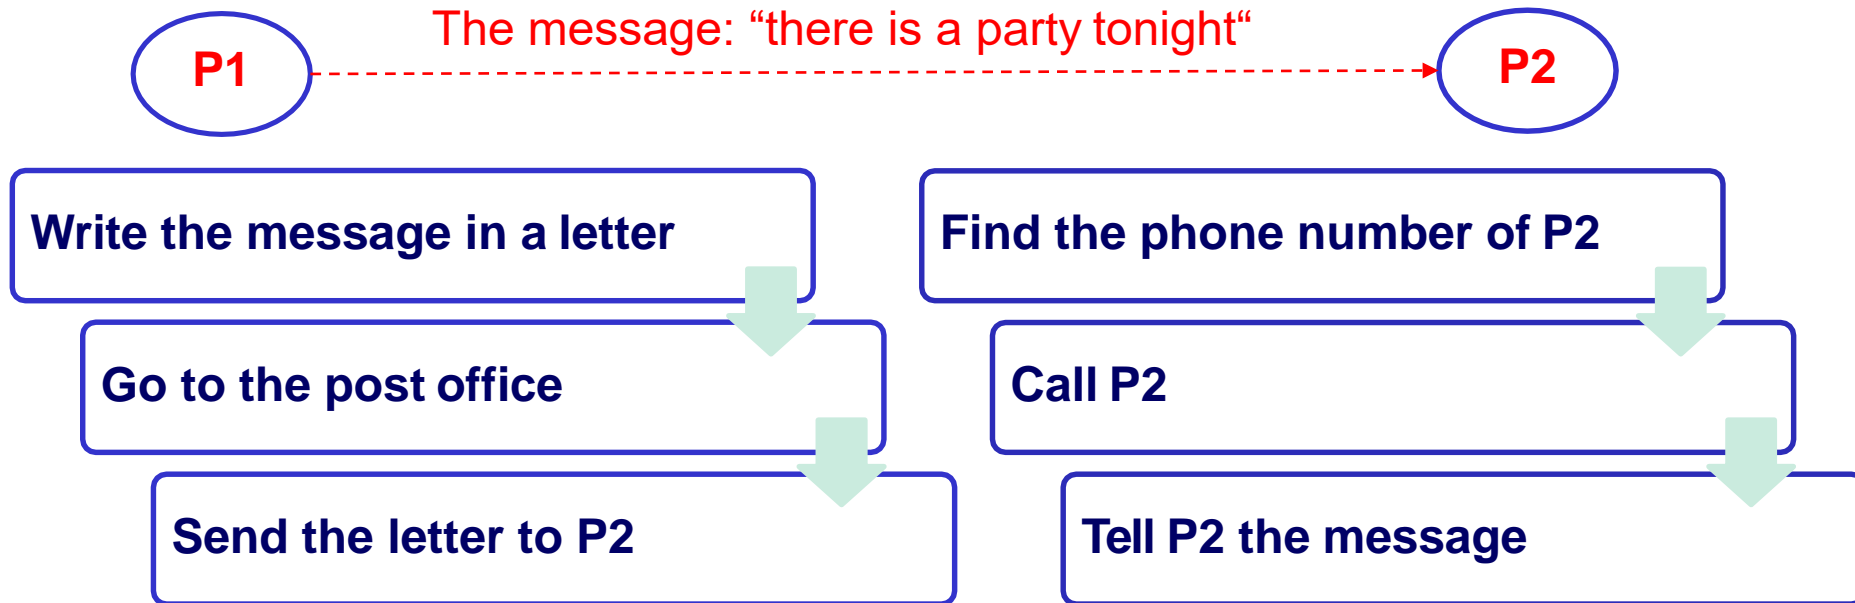
```
$sudo apt-get install mpich  
$mpicc c_ex04.c  
$mpirun -np 4  
./a.out
```

```
MPI_Comm_size(MPI_COMM_WORLD,&numprocs);  
MPI_Comm_rank(MPI_COMM_WORLD,&myid);  
source=0;  
count=4;  
if(myid == source){  
    for(i=0;i<count;i+  
        +) buffer[i]=i;  
}  
MPI_Bcast(buffer,count,MPI_INT,source,MPI_COMM_WORLD);
```

[http://geco.mines.edu/workshop/class2/examples/mpl/c\\_ex04.c](http://geco.mines.edu/workshop/class2/examples/mpl/c_ex04.c)




# Connection-oriented or connectionless communication



Connection-oriented communication between P1 and P2 requires the setup of communication connection between them first – no setup in connectionless communication

# Blocking versus non-blocking communication calls



**Send:** transmitting a message is finished, it does not necessarily mean that the message reaches its final destination.

- **Blocking:** the process/thread execution is suspended until the message transmission finishes
- **Non-blocking:** the process/thread execution continues without waiting until the finish of the message transmission

# Persistent and transient communication

---


- Persistent communication
  - Messages are kept in the communication system until they are delivered to the receiver
  - Often storage is needed
- Transient communication
  - Messages are kept in the communication temporary only if both the sender and receiver are live

# Asynchronous versus synchronous communication

---



- Asynchronous: the process continues after as soon as sending messages have been copied to the local buffer
  - Non blocking send; receive may/may not be blocking
  - Callback mechanisms
- Synchronous: the sender **waits until it knows** the messages have been **delivered** to the receiver
  - Blocking send/blocking receive
  - Typically utilize connection-oriented and keep-alive connection
  - Blocking request-reply styles

# Different forms of communication




Source: Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems – Principles and Paradigms*, 2nd Edition, 2007, Prentice-Hall

# Stateful versus Stateless Server

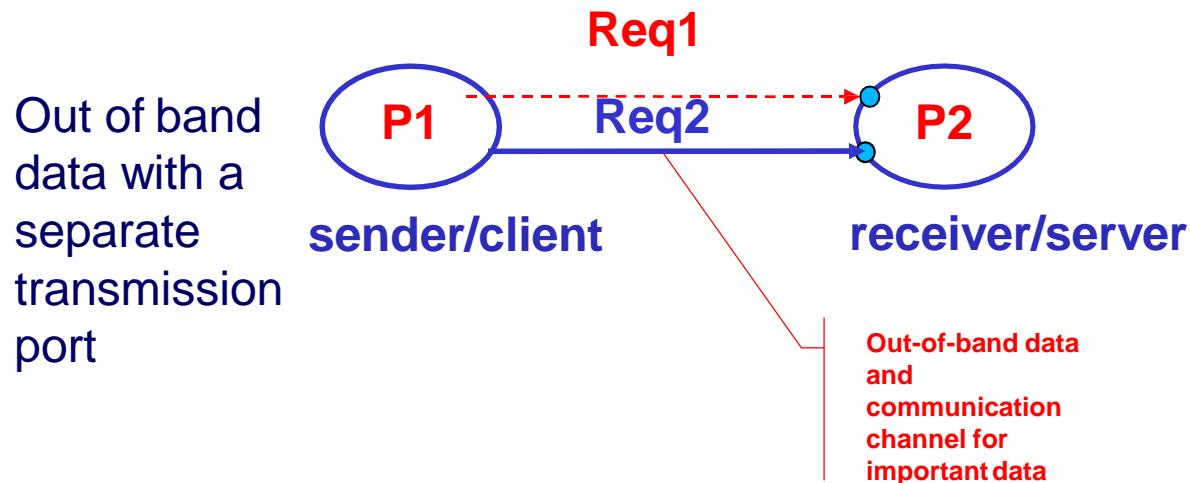


Stateless server	Soft State	Stateful Server
Does not keep client's state information	Keep some limited client's state information in a limited time	Maintain client's state information permanently

# Handling of band data



All messages come to P2 in the same port, no clear information about priority

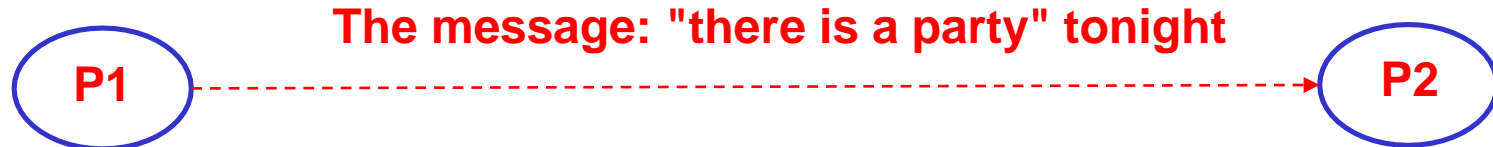


---

# **COMMUNICATION PROTOCOLS**



# Some key questions - Protocols




- **Communication patterns**
  - Can I use a single sending command to send the message to multiple people?
- **Identifier/Naming/Destination**
  - How do I identify the guys I need to send the message
- **Connection setup**
  - Can I send the message without setting up the connection
- **Message structure**
  - Can I use German or English to write the message
- **Layered communication**
  - Do I need other intermediators to relay the message?
- ...


**A communication protocol will describe rules addressing these issues**

# Applications and Protocols

## Application-specific protocols



## Application-independent protocols



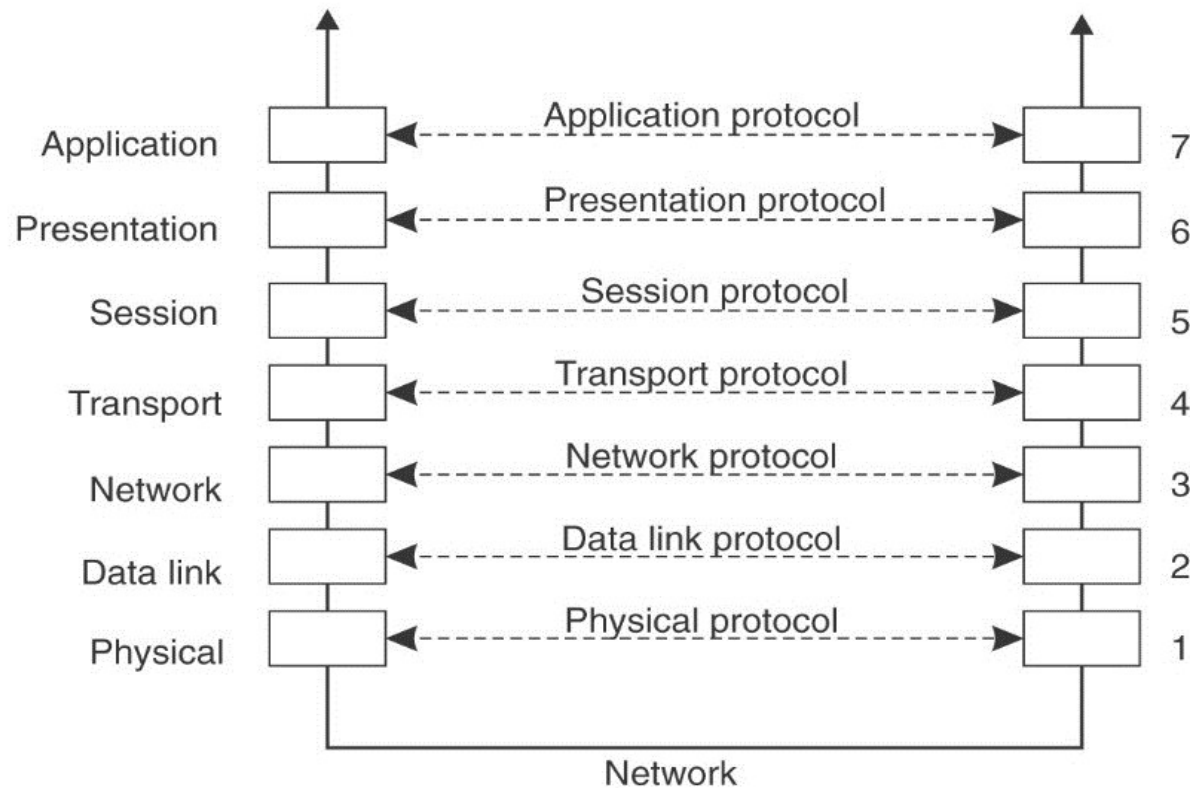
Source: Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems – Principles and Paradigms*, 2nd Edition, 2007, Prentice-Hall

# Layered Communication Protocols

---








- Complex and open communication requires **multiple communication protocols**
- Communication protocols are typically organized into different layers: layered protocols/protocol stacks
- Conceptually: each layer has a set of different **protocols for certain communication functions**
  - Different protocols are designed for different environments/criteria
- **A protocol suite**: usually a set of protocols used together in a layered model

# OSI – Open Systems Interconnection Reference Model



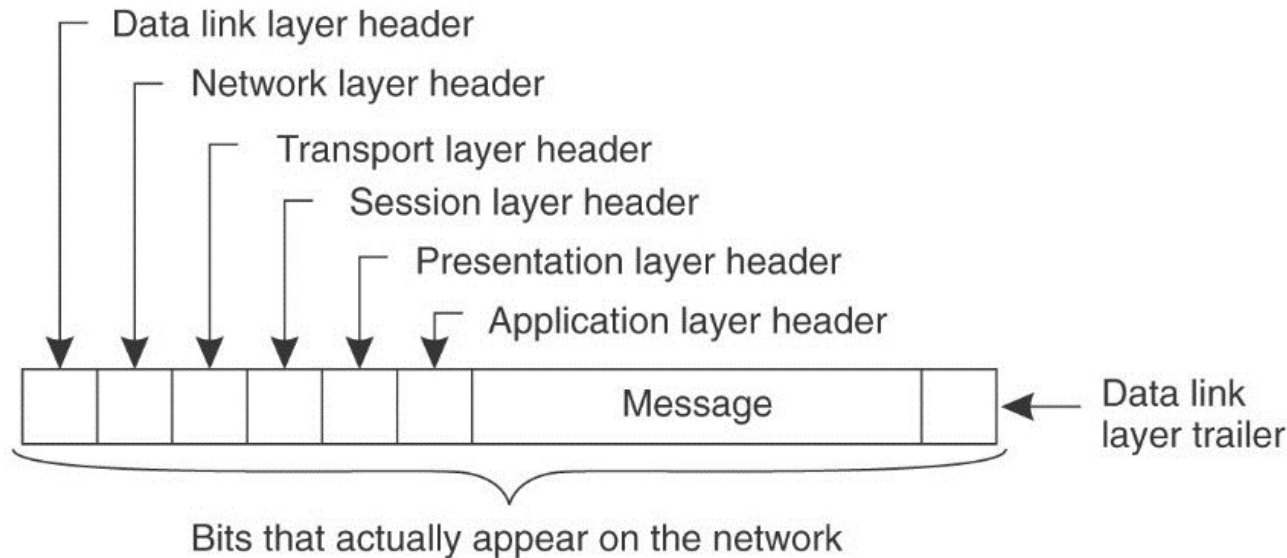
Source: Andrew S. Tanenbaum and Maarten van Steen, Distributed Systems – Principles and Paradigms, 2nd Edition, 2007, Prentice-Hall

# OSI Layers

Application		<ul style="list-style-type: none"><li>• Support application-specific needs</li></ul>
Presentation		<ul style="list-style-type: none"><li>• Process information format and deliver the information for the application layer (e.g., serializing and encryption)</li></ul>
Session		<ul style="list-style-type: none"><li>• Manage communication sessions between applications</li></ul>
Transport		<ul style="list-style-type: none"><li>• Provide an end-to-end communication for applications by delivering data among applications</li></ul>
Network		<ul style="list-style-type: none"><li>• Route data packets among senders/receivers</li></ul>
Data Link		<ul style="list-style-type: none"><li>• Deal with sending data frames (units of bits) and detecting/correcting data frames</li></ul>
Physical Layer		<ul style="list-style-type: none"><li>• Transfer binary data (bits) over physical interfaces (e.g., fiber optics)</li></ul>

# How layered protocols work – message exchange

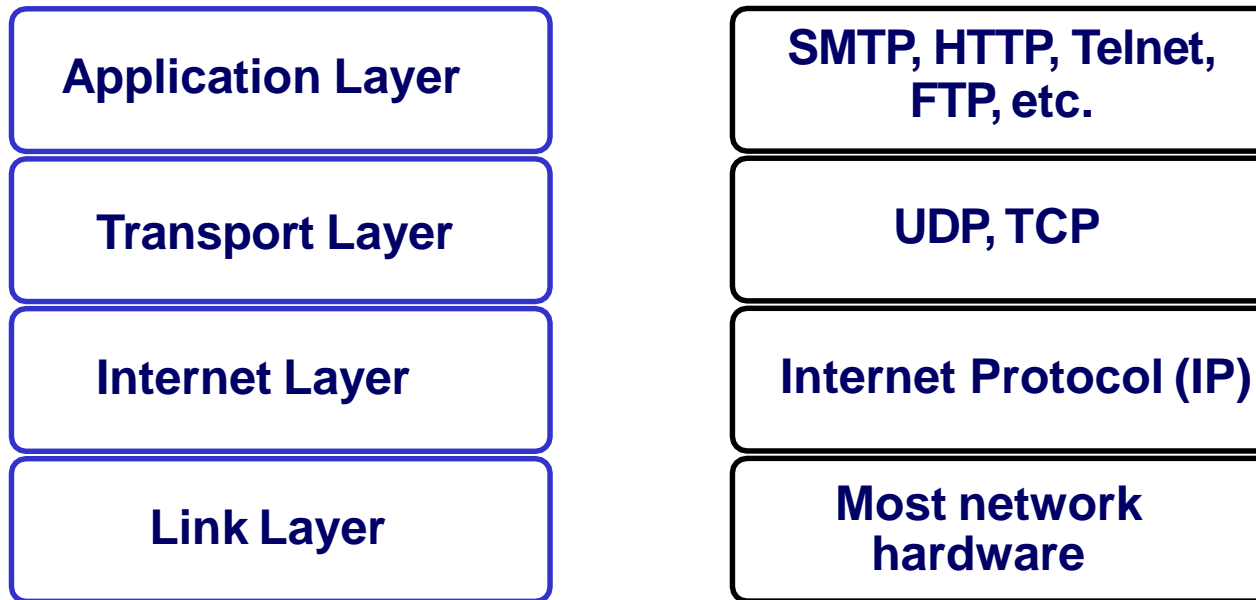
- Principles of constructing messages/data encapsulation



Source: Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems – Principles and Paradigms*, 2nd Edition, 2007, Prentice-Hall

# TCP/IP

- The most popular protocol suite used in the Internet
- Four layers



<http://tools.ietf.org/html/rfc1122>

# Internet Protocol (IP)

- Defines the datagram as the basic data unit
- Defines the Internet address scheme
- Transmits data between the Network Access Layer and Transport Layer
- Routes datagrams to destinations
- Divides and assembles datagrams




Figure source:

[http://en.wikipedia.org/wiki/Internet\\_protocol\\_suite](http://en.wikipedia.org/wiki/Internet_protocol_suite)



# TCP/IP – Transport Layer

- Host-to-host transport features
- Two main protocols: TCP (Transmission Control Protocol) and UDP (User Datagram Protocol)

Layer\Protocol	TCP	UDP
Application layer	Data sent via Streams of bytes	Data sent in Messages
Transport Layer	Segment	Packet
Internet Layer	Datagram	Datagram
Link Layer	Frame	Frame

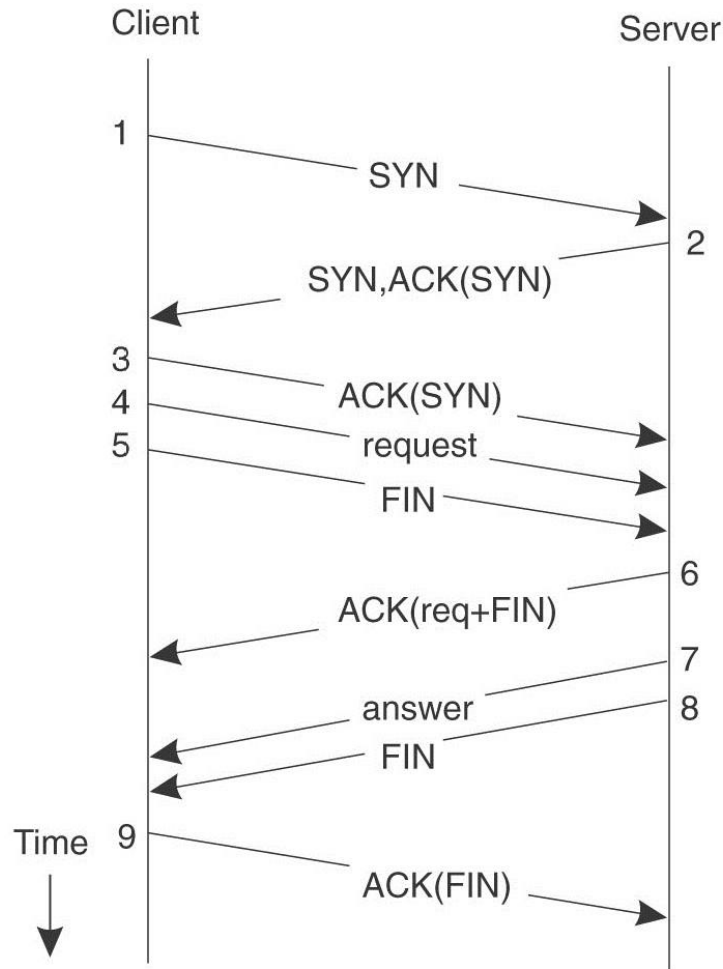
**Segment is the original data + Transport Layer header.**

**Packet is a Segment + Network Layer header.**

**Frame is a Packet + Data Link Layer header.**

**Note: pay attention with the terms, packet/datagram“ in TCP/IP versus that in the OSI model**

# TCP Operations



**SYN: Synchronize**  
TCP packet sent to server  
requesting that a connection  
be established

**ACK(SYN)**

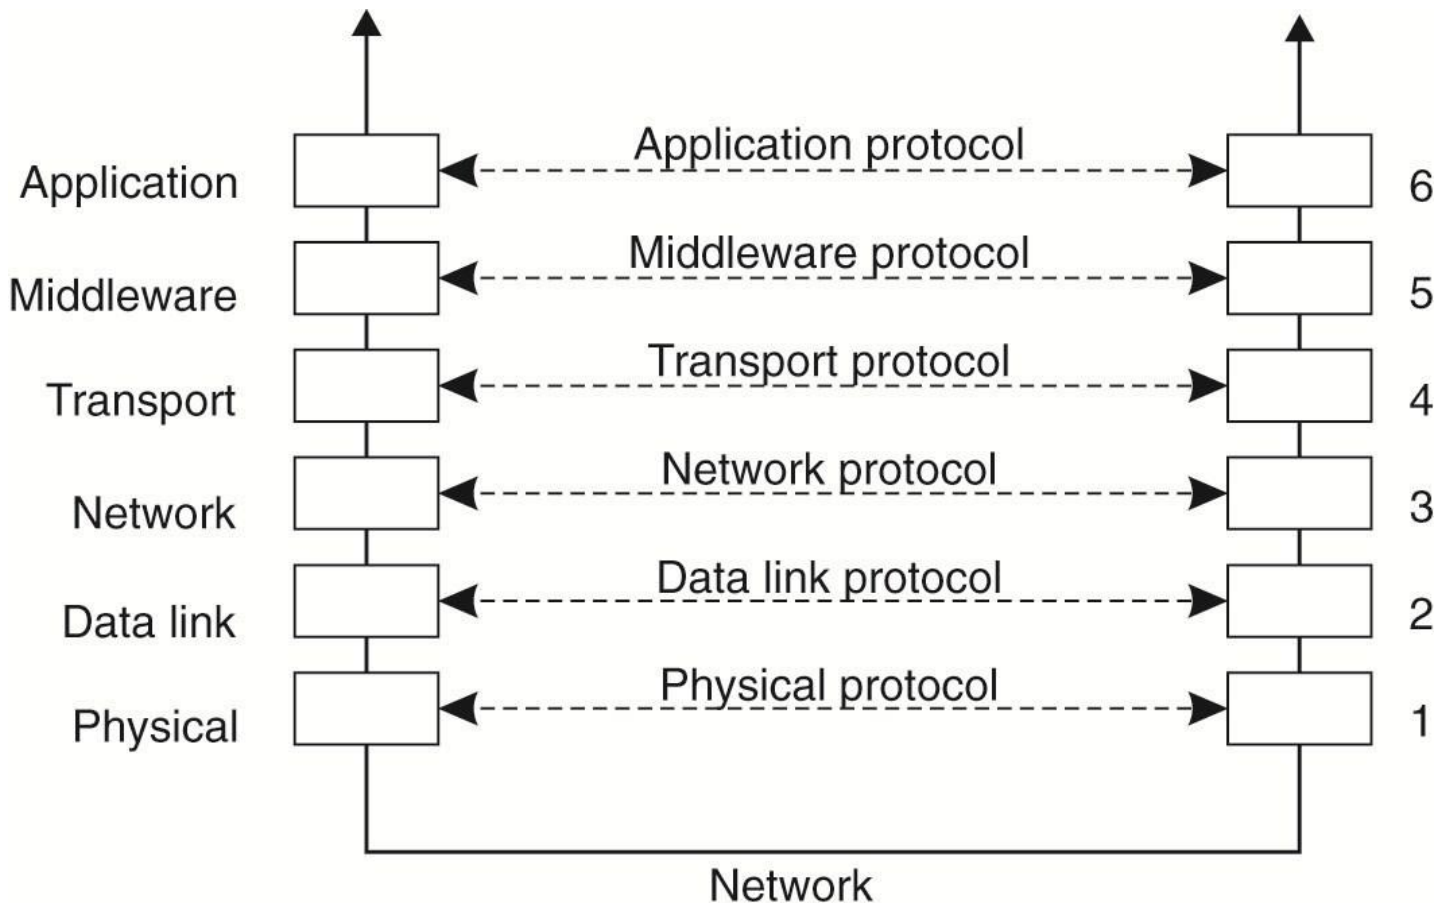
If the SYN is received by the  
server, an SYN/ACK is sent  
back to the address  
requested by the SYN

# Communication protocols are not enough

---

- We need more than just communication protocols
  - E.g., resolving names, electing a communication coordinator, locking resources, and synchronizing time
- Middleware
  - Including a set of general-purpose but application-specific protocols, middleware communication protocols, and other specific services.

# Middleware Protocols



Source: Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems – Principles and Paradigms*, 2nd Edition, 2007, Prentice-Hall

---

# **HANDLING COMMUNICATION MESSAGES/REQUESTS**

# Where communication takes place?

---

- Message passing – send/receive
  - Processes **send** and **receive** messages
    - Sending process versus receiving process
    - Communication is done by using **a set of functions for communication implementing protocols**
- Remote method/procedure calls
  - A process **calls/invokes a (remote) procedure** in another process
    - Local versus remote procedure call, but in the same manner
- Remote object calls
  - A process **calls/invokes a (remote) object** in another process

# Basic send/receive communication

```
# Echo client program
import socket
```

```
HOST = 'daring.cwi.nl' # The remote host
PORT = 50007           # The same port as
                        # used by the server
s = socket.socket(socket.AF_INET,
                  socket.SOCK_STREAM)
s.connect((HOST, PORT))
s.send('Hello, world')
data = s.recv(1024)
s.close()
print 'Received', repr(data)
```

Network



```
# Echo server program
import socket
```

```
HOST = "                # Symbolic name meaning the
                        # local host
PORT = 50007            # Arbitrary non-privileged
                        # port
s = socket.socket(socket.AF_INET,
                  socket.SOCK_STREAM)
s.bind((HOST, PORT))
s.listen(1)
conn, addr = s.accept()
print 'Connected by', addr
while 1:
    data = conn.recv(1024)
    if not data: break
    conn.send(data)
conn.close()
```

Python source: <http://docs.python.org/release/2.5.2/lib/socket-example.html>

# Remote procedure calls (RPC)

```
void
hello_prog_1(char *host)
{
    CLIENT *clnt;
    char **result_1;
    char *hello_1_arg;

#ifdef DEBUG
    clnt = clnt_create (host, HELLO_PROG, HELLO_VERS, "udp");
    if (clnt == NULL) {
        clnt_pcreateerror (host);
        exit(1);
    }
#endif /* DEBUG */

    result_1 = hello_1((void*)&hello_1_arg, clnt);
    if (result_1 == (char **) NULL) {
        clnt_perror (clnt, "call failed");
    }
#ifdef DEBUG
    clnt_destroy (clnt);
#endif /* DEBUG */
    printf("result is: %s\n",(*result_1));
}

int
main (int argc, char *argv[])
{
    char *host;

    if (argc < 2) {
        printf ("usage: %s server_host\n", argv[0]);
        exit(1);
    }
    host = argv[1];
    hello_prog_1 (host);
    exit(0);
}
```

Network

## Procedure in a remote server

```
char **
hello_1_svc(void *argp, struct svc_req *rqstp)
{
    static char * result ="Hello";

    /*
     * insert server code here
     */

    return &result;
}
```



# Remote procedure calls (RPC)

## Objects in a remote server

```
public class ComputePi {
    public static void main(String args[]) {
        if (System.getSecurityManager() == null) {
            System.setSecurityManager(new SecurityManager());
        }
        try {
            String name = "Compute";
            Registry registry = LocateRegistry.getRegistry(args[0]);
            Compute comp = (Compute) registry.lookup(name);
            Pi task = new Pi(Integer.parseInt(args[1]));
            BigDecimal pi = comp.executeTask(task);
            System.out.println(pi);
        } catch (Exception e) {
            System.err.println("ComputePi exception:");
            e.printStackTrace();
        }
    }
}
```

```
public interface Compute extends Remote {
    <T> T executeTask(Task<T> t) throws RemoteException;
}
....
public class ComputeEngine implements Compute {
    public ComputeEngine() {
        super();
    }
    public <T> T executeTask(Task<T> t) {
        return t.execute();
    }
    public static void main(String[] args) {
        if (System.getSecurityManager() == null) {
            System.setSecurityManager(new SecurityManager());
        }
        try {
            String name = "Compute";
            Compute engine = new ComputeEngine();
            Compute stub =
                (Compute) UnicastRemoteObject.exportObject(engine, 0);
            Registry registry = LocateRegistry.getRegistry();
            registry.rebind(name, stub);
            System.out.println("ComputeEngine bound");
        } catch (Exception e) {
            System.err.println("ComputeEngine exception:");
            e.printStackTrace();
        }
    }
}
```

### Java Source:

<https://docs.oracle.com/javase/tutorial/rmi/overview.html>


# Processing multiple requests

---

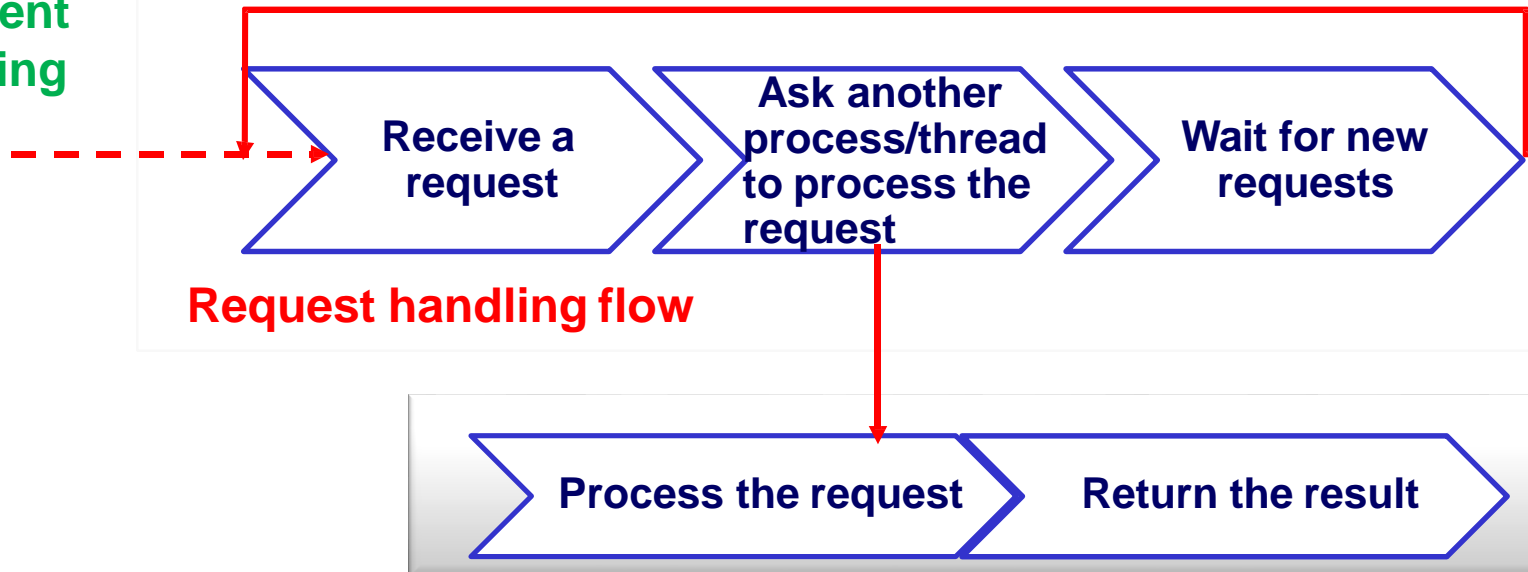
- How to deal with multiple, concurrent messages received?
- Problems:
  - Different roles: clients versus servers/services
    - A large number of clients interact with a small number of servers/services
    - A single process might receive a lot of messages at the same time
- Impacts
  - performance, reliability, cost, etc.

# Iterative vs Concurrent processing

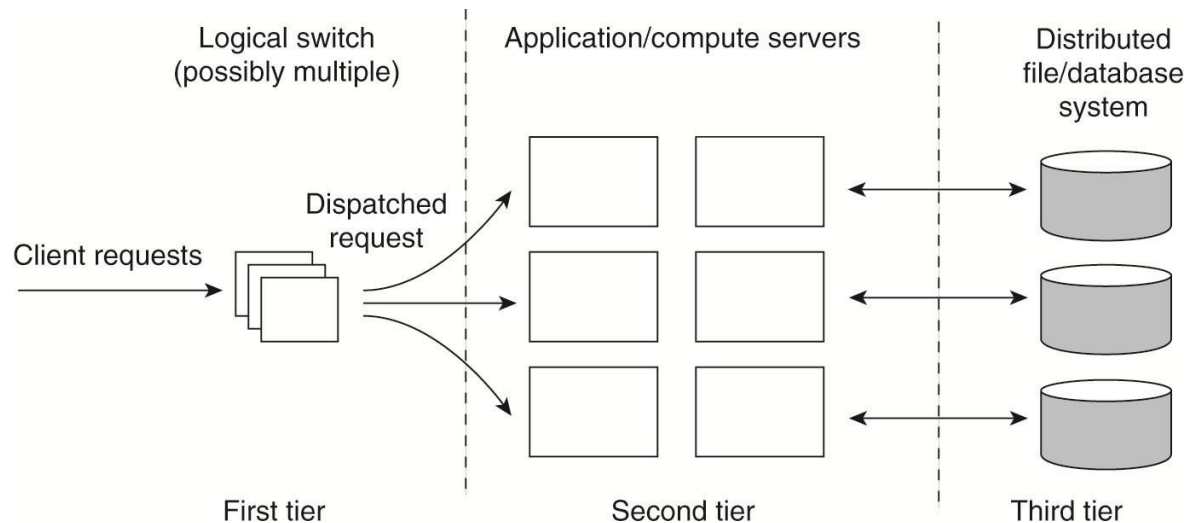
Iterative  
processing



Concurrent  
processing



# Using replicated processes

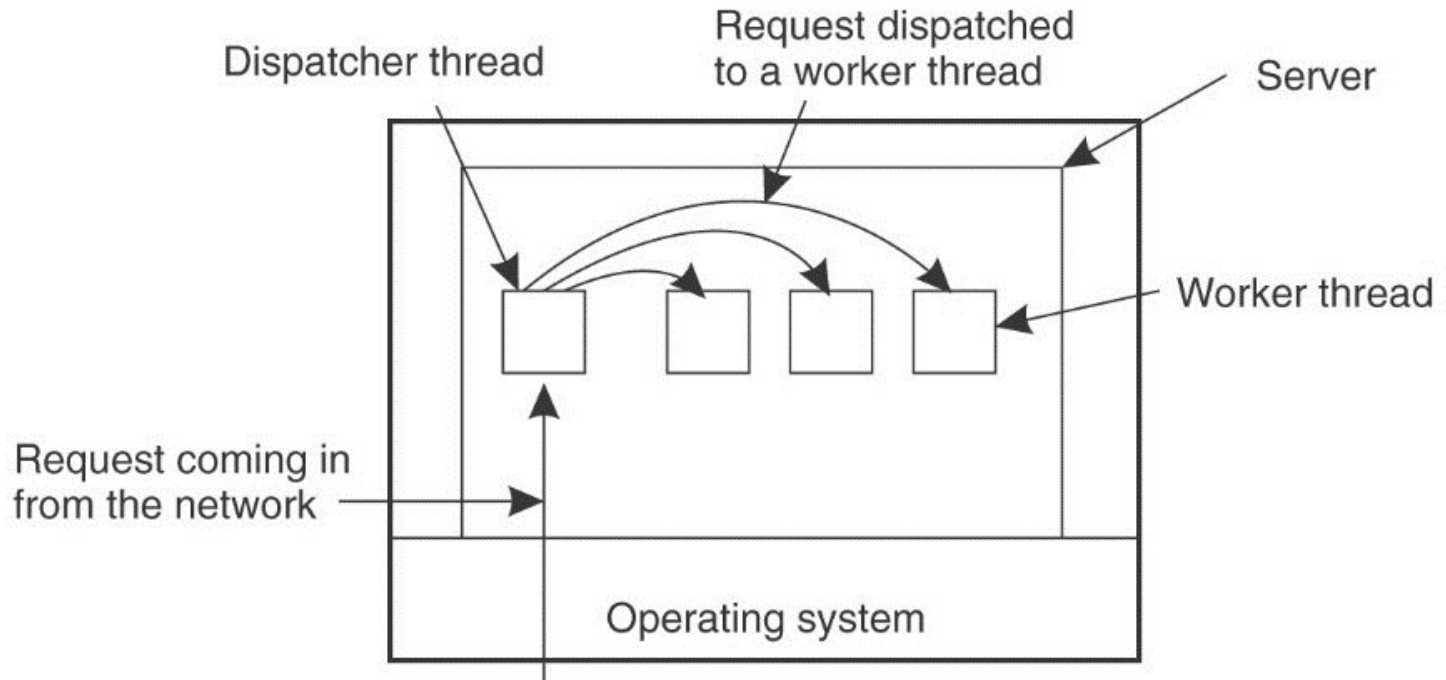


**Often load  
balancing  
mechanisms  
are needed**

Source: Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems – Principles and Paradigms*, 2nd Edition, 2007, Prentice-Hall

**Q: How does this model help to improve performance and fault-tolerance? What would be a possible mechanism to reduce costs based on the number of client requests?**

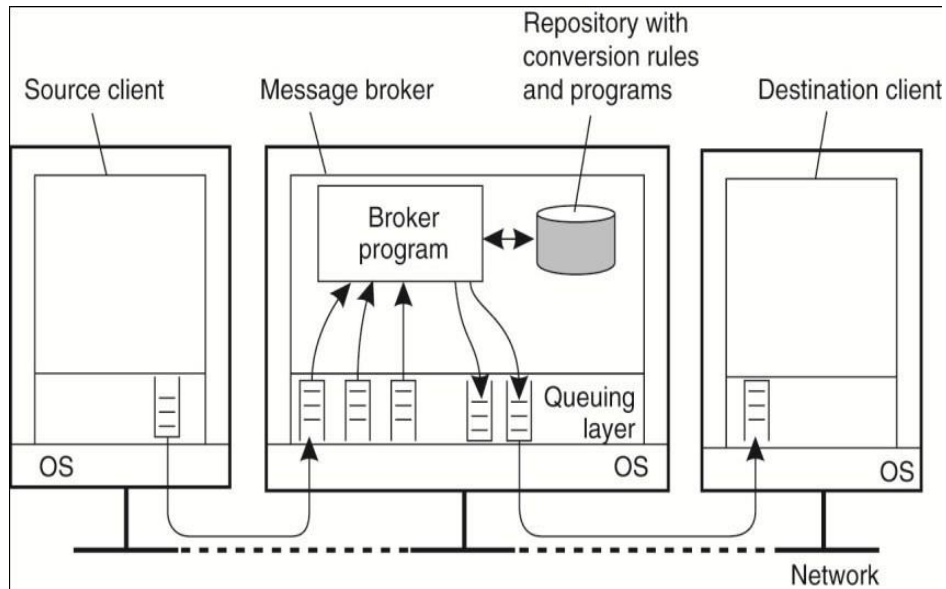
# Using multiple threads




Source: Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems – Principles and Paradigms*, 2nd Edition, 2007, Prentice-Hall

Q: Compare this architectural model with the super-server model?

# Using message brokers/space repository



Source: Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems – Principles and Paradigms*, 2nd Edition, 2007, Prentice-Hall



# Example

- Get a free instance of RabbitMQ from cloudamqp.com
- Get code from: <https://github.com/cloudamqp/java-amqp-example>
- First run the test sender, then run the receiver



# Summary

---

- Complex and diverse communication patterns, protocols and processing models
- Choices are based on communication requirements and underlying networks
  - Understand their pros/cons
  - Understand pros and cons of their technological implementations
- Dont forget to play with some simple examples to understand existing concepts